



Northeastern

EECE5698

Parallel Processing for Data Analytics

Lecture 3: Map and Reduce Operations in Spark

Outline

- ❑ Map and other Transforms
- ❑ Reduce and other Actions

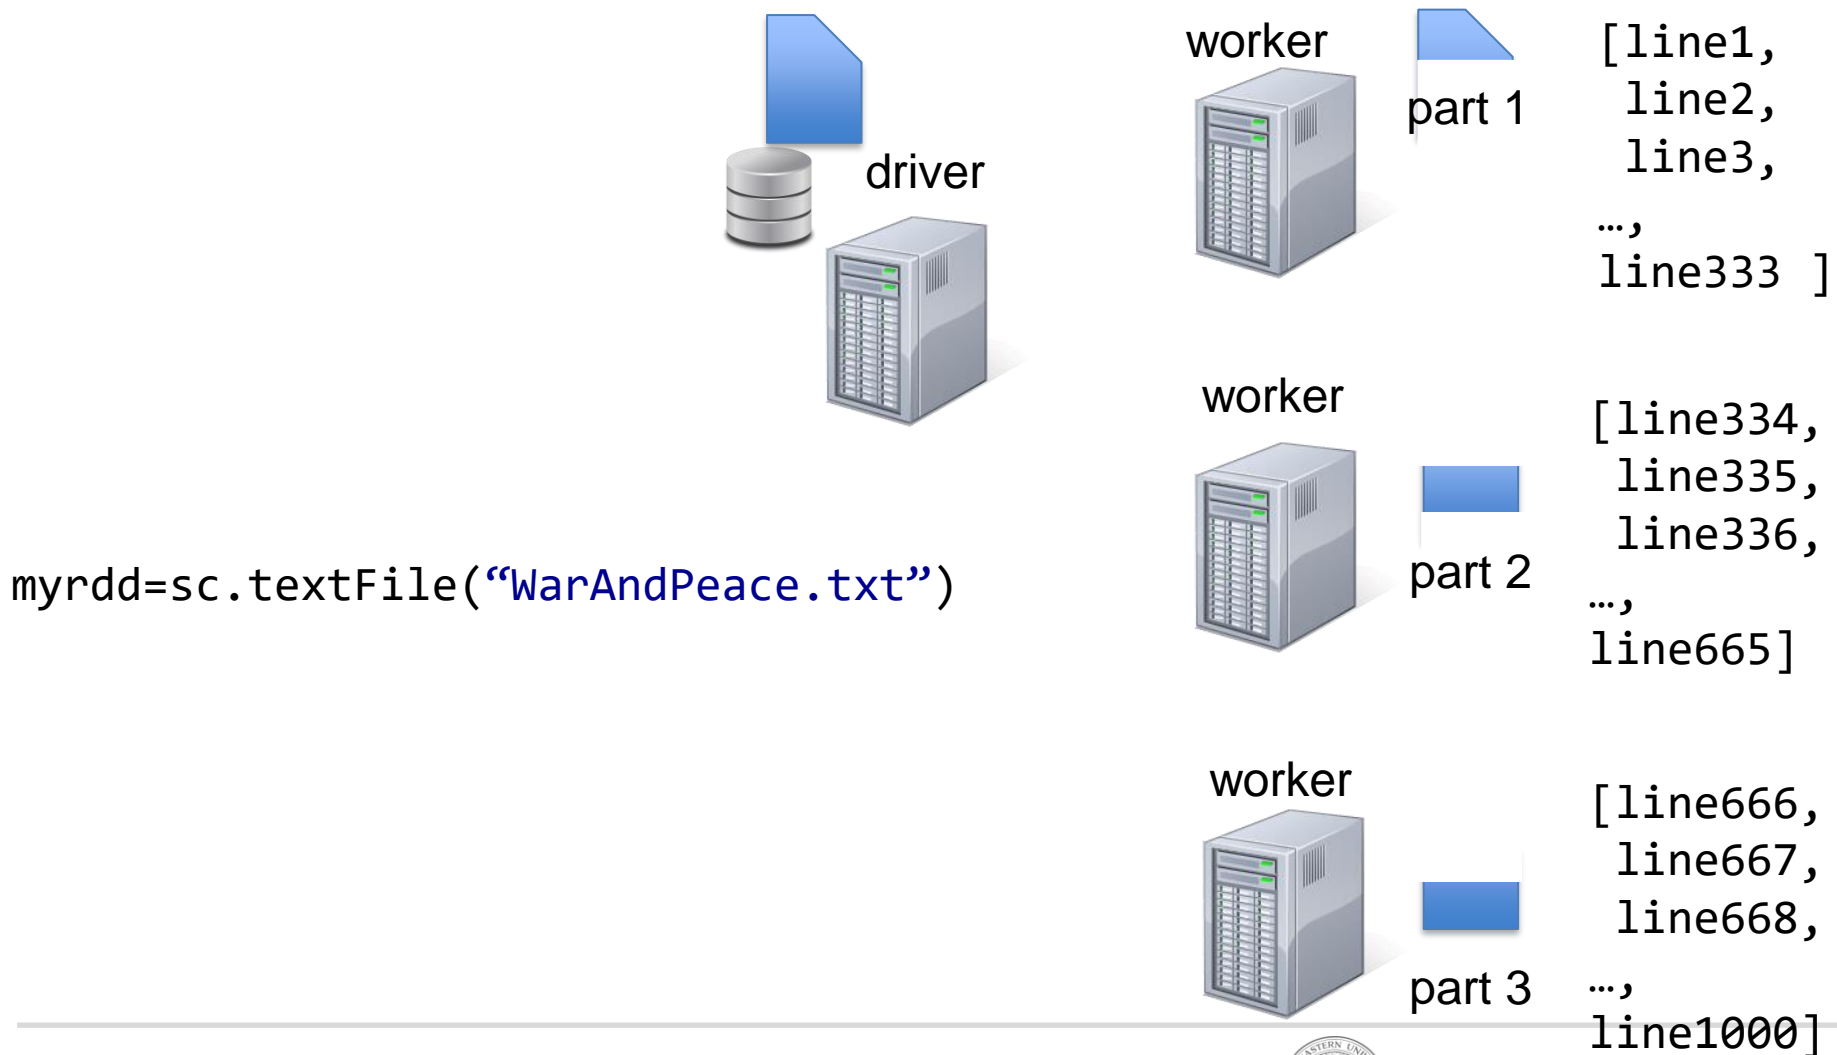


Outline

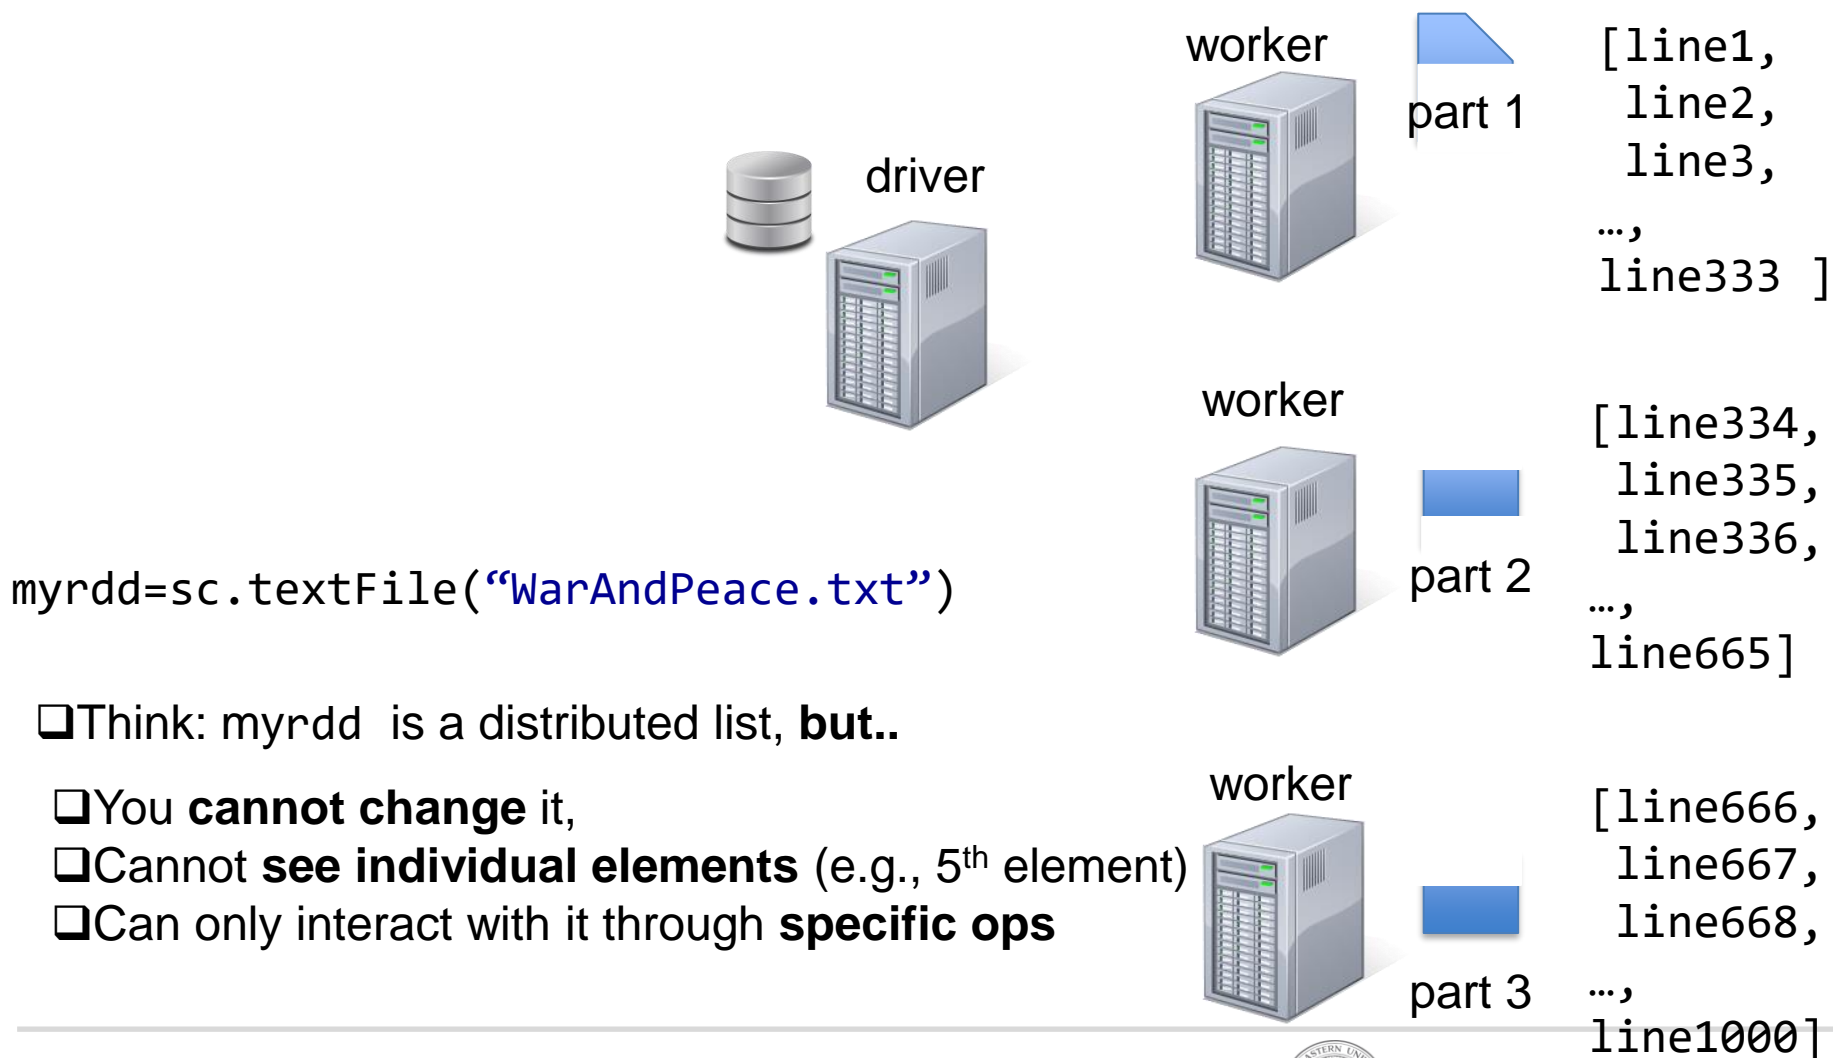
- ❑ Map and other Transforms
- ❑ Reduce, and other Actions



Resilient Distributed Datasets (RDDs)



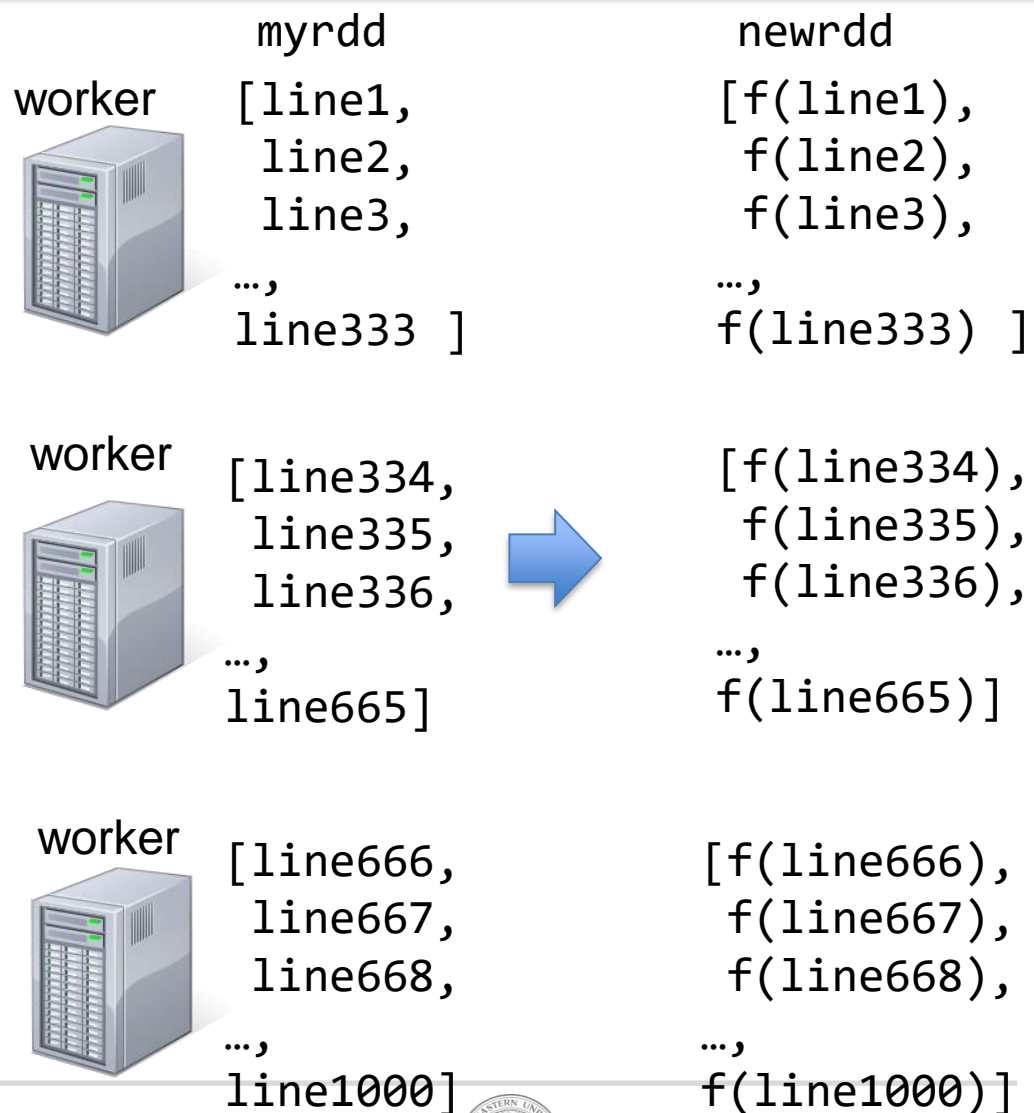
Resilient Distributed Datasets (RDDs)



The Map Transformation

```
newrdd=myrdd.map(f)
```

- Return a new RDD by applying a function to each element of this RDD.



Examples

User-defined function, defined through **def** statement

```
def myf(x):  
    return x+2
```

```
newrdd=myrdd.map(myf)
```

myrdd

[1,
2,
3,
...,
10]



newrdd

[3,
4,
5,
...,
12]

Examples

User-defined function, defined through **lambda** operator

```
newrdd=myrdd.map(lambda x:x+2)
```

myrdd

[1,
2,
3,
...,
10]



newrdd

[3,
4,
5,
...,
12]

Examples

User-defined function, defined through **lambda** operator

```
newrdd=myrdd.map(lambda x:x+'_')
```

myrdd

```
[ 'a',  
  'b',  
  'c',  
  ...,  
  'z' ]
```



newrdd

```
[ 'a_',  
  'b_',  
  'c_',  
  ...,  
  'z_' ]
```

Examples

Python built-in function

```
newrdd=myrdd.map(eval)
```

myrdd

```
[ '1.4',  
  '2.3',  
  '(1,3)',  
  '{“apples”:2,“oranges”:3}',  
  'array([1,3,5])'  
]
```



newrdd

```
[1.4,  
 2.3,  
(1,3),  
{“apples”:2,“oranges”:3},  
array([1.0,3.0,5.0])  
]
```

Other Important Transforms

`flatMap(f)`: Map each element to multiple elements

```
newrdd=myrdd.flatMap(lambda x:x.split())
```

myrdd

```
['Words are like leaves',  
'and where most abound',  
'much fruit of sense beneath',  
'is most rarely found']
```



newrdd

```
['Words',  
'are',  
'like',  
'leaves',  
'and',  
'where',  
...,  
'rarely',  
'found']
```

NOTE: `f` must return a **collection** (e.g., list)

Other Important Transforms

`distinct()`: Remove duplicates

```
newrdd= myrdd.distinct()
```

myrdd

```
[1,  
 1,  
 2,  
 2,  
 3  
]
```



newrdd

```
[1,  
 2,  
 3  
]
```

Other Important Transforms

`sample(withReplacement, fraction)`: Return a random sampled subset of this RDD

```
newrdd= myrdd.sample(False,0.5)
```

myrdd

```
[1,  
 2,  
 3,  
 4,  
 5  
]
```



newrdd

```
[1,  
 3,  
 5  
]
```

(result is non-deterministic)

Other Important Transforms

`sortBy(keyfunc, ascending=True)`: Sort RDD contents

```
newrdd= myrdd.sortBy(lambda x:-x)
```

myrdd

```
[1,  
 4,  
 5,  
 2,  
 3  
]
```



newrdd

```
[5,  
 4,  
 3,  
 2,  
 1  
]
```

Warning: future transforms may change ordering

How does map work?

```
def myf(x):  
    return x+2
```



1

010
101

driver



Serialize to
binary
representation

worker



[line1,
line2,
line3,
...,
line333]

worker



[line334,
line335,
line336,
...,
line665]

worker



[line666,
line667,
line668,
...,
line1000]

How does map work?

```
def myf(x):  
    return x+2
```



1

010
101

Serialize to
binary
representation



2 Ship to
workers



010
101

worker



[line1,
line2,
line3,
...,
line333]



010
101

worker



[line334,
line335,
line336,
...,
line665]



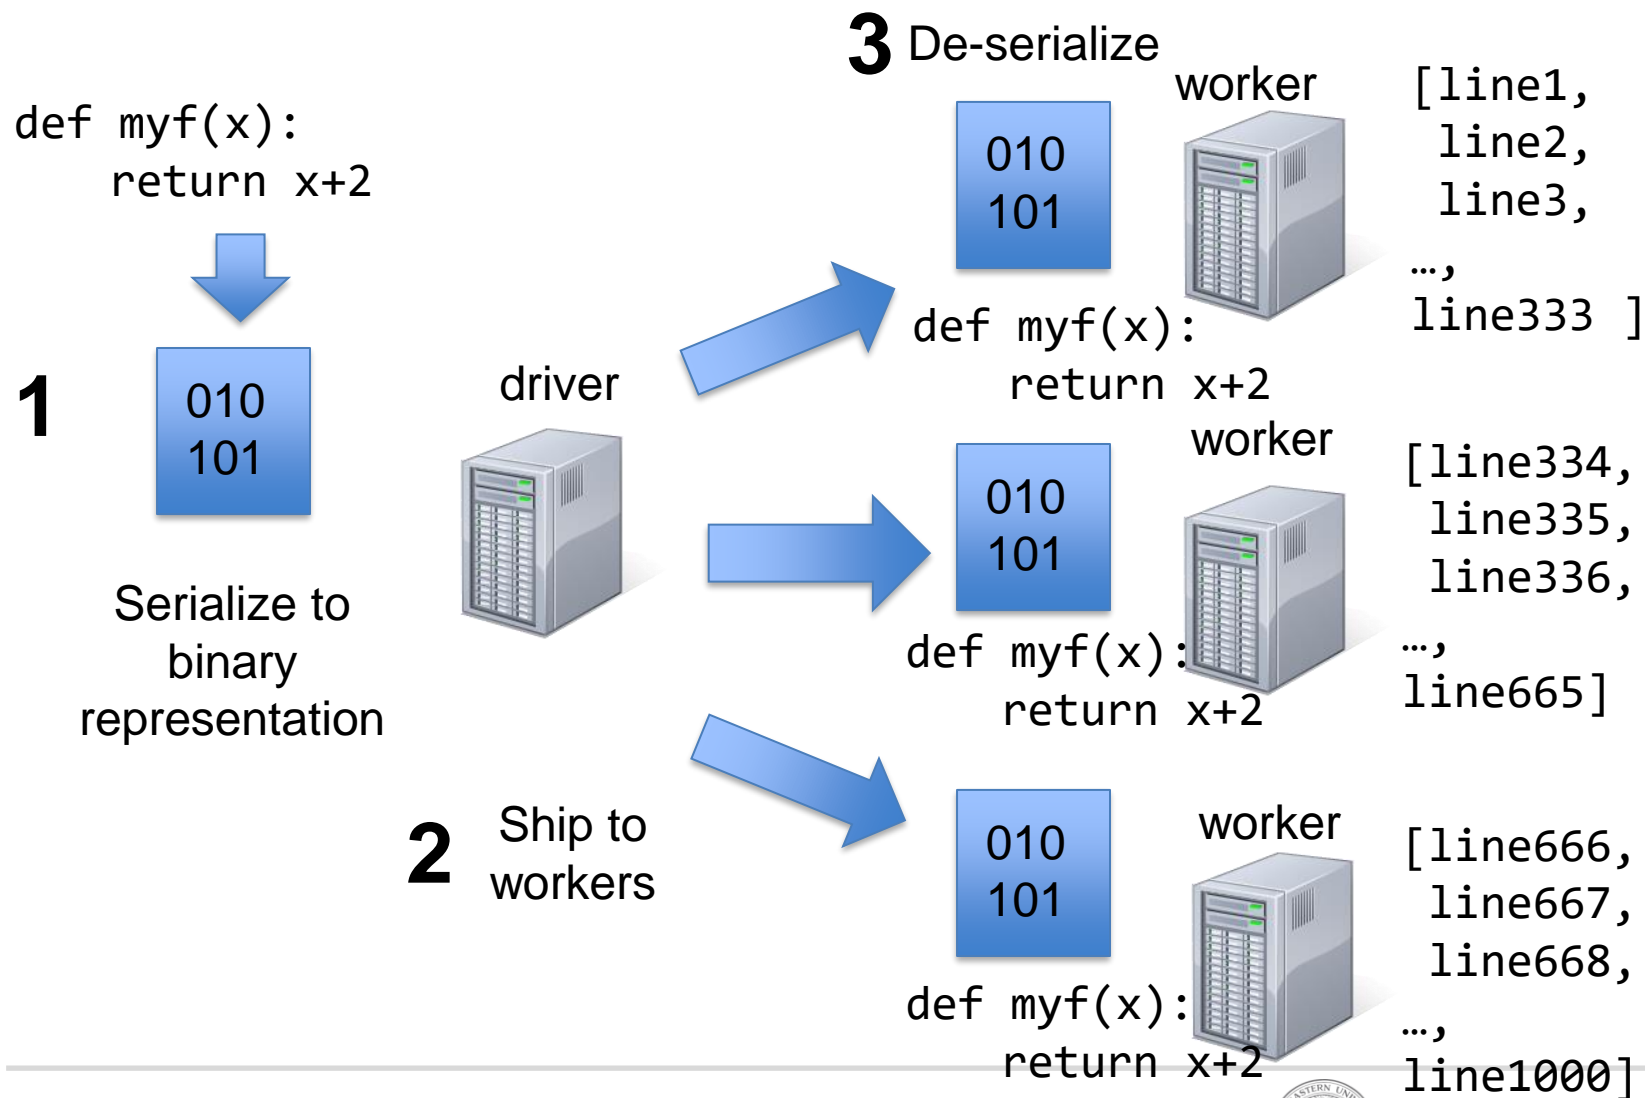
010
101

worker

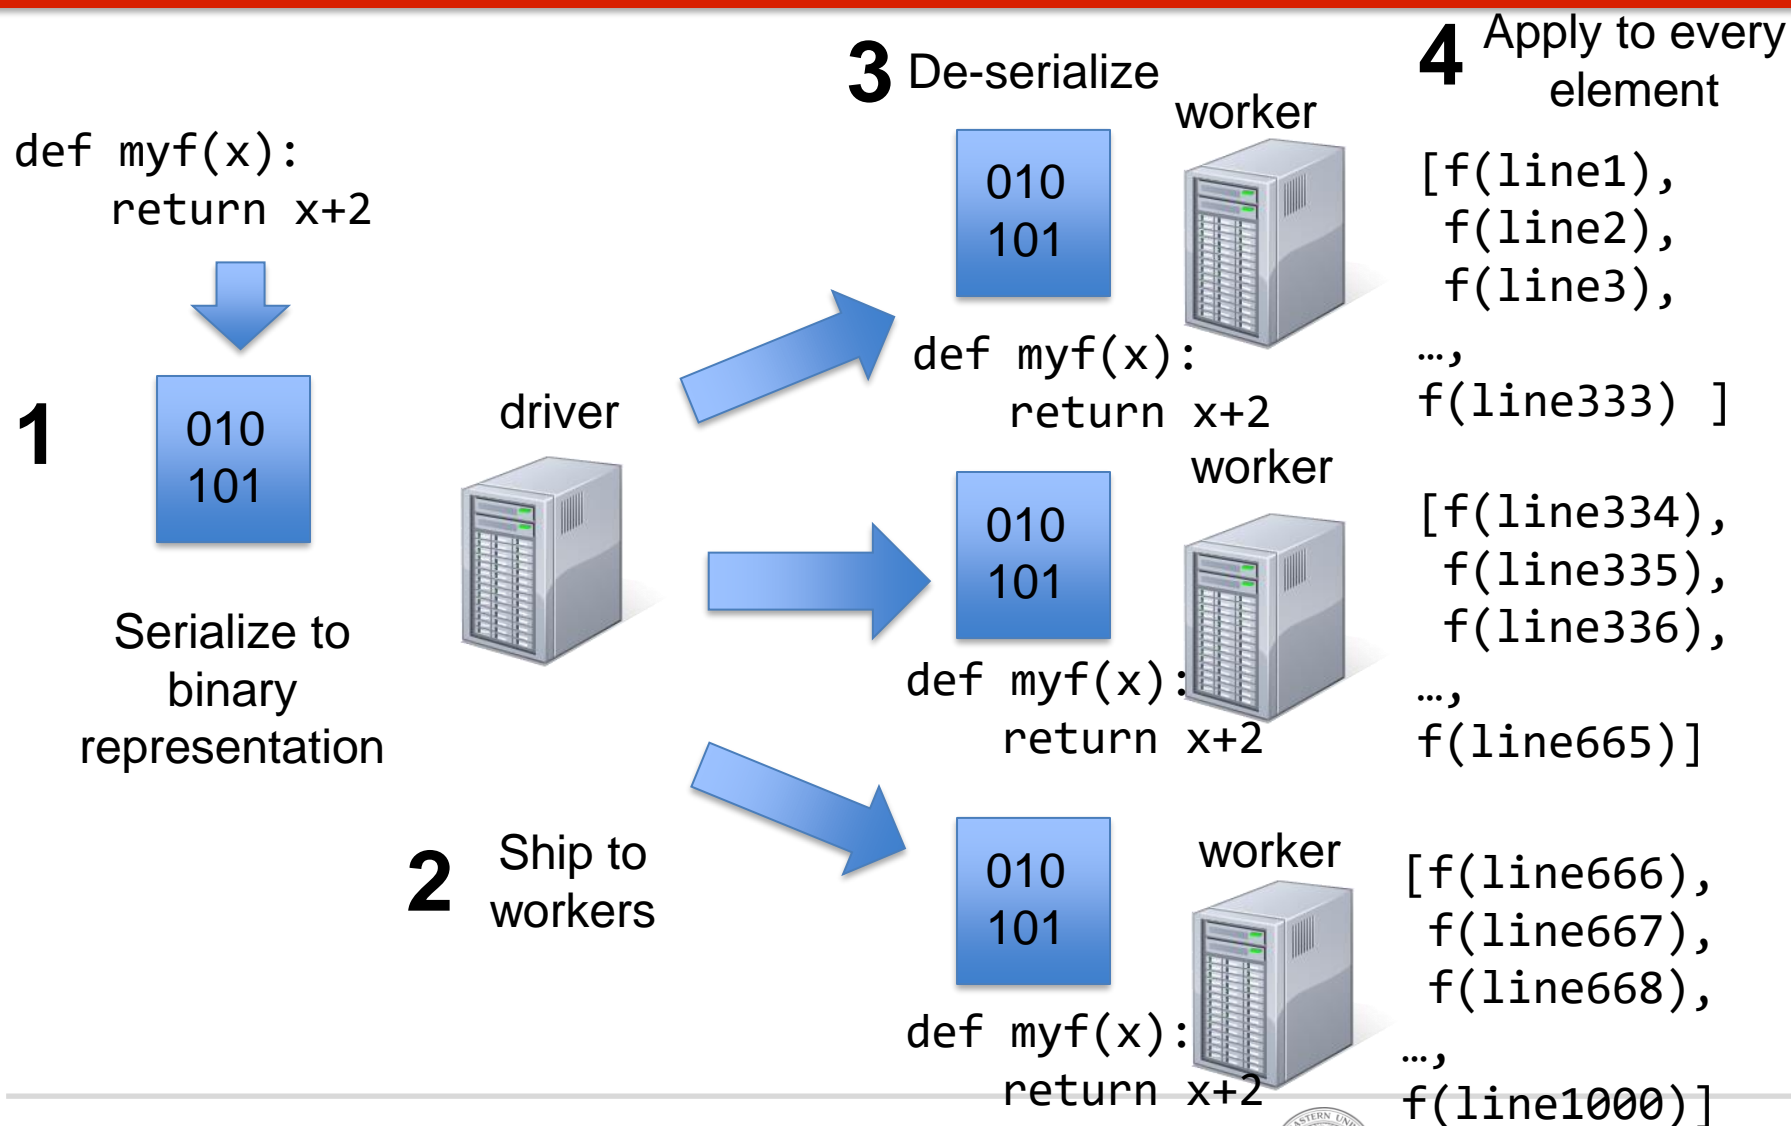


[line666,
line667,
line668,
...,
line1000]

How does map work?



How does map work?



maps Involve Communication

- ❑ **Function f is shipped** to data
- ❑ **Data do not move** (stay at workers)
- ❑ **Serialize-deserialize:** `pickle` (default option)



Shipping Local Variables

- ❑ **Variables defined in the driver program** will automatically be shipped to the cluster along with function definition:

```
query = raw_input("Enter a query:")  
pages.filter(lambda x: x.startswith(query))
```

- ❑ Some caveats:
 - Each task gets a new copy (**updates aren't sent back**)
 - Variable **must be pickle-able**
 - Variable **cannot be** an rdd
 - Don't use **fields of an object** (ships all of it!)
 - Beware of **shipping large variables!**



Example

```
class myClass(object):  
    def __init__(self):  
        self.mylist = range(10000) # 10000 element list  
        self.query = "ERROR"
```

```
myOb=myClass()
```

```
pages.filter(lambda x: x.startswith(myOb.query))
```

**Ships ENTIRE myOB , including mylist
(no error reported)**



Example

```
class myClass(object):  
    def __init__(self):  
        self.myfile = file('temp.pkl', 'w') # file handle  
        self.query = "ERROR"  
  
myOb=myClass()  
  
pages.filter(lambda x: x.startswith(myOb.query))
```

This will produce an error (attempts to ship myObj, which contains non-picklable myfile)



Example

```
class myClass(object):  
    def __init__(self):  
        self.myrdd = sc.textFile('WarAndPeace.txt') # rdd  
        self.query = "ERROR"  
  
myOb=myClass()  
  
pages.filter(lambda x: x.startswith(myOb.query))
```

This will produce an error (attempts to ship myObj, and rdds are not allowed in functions)




Outline

- Map and other Transforms
- Reduce and other Actions




The Reduce Action

```
val=myrdd.reduce(f)
```

worker  myrdd
[x1,
x2,
x3,
...,
x333]

```
def add(x,y):  
    return x+y
```

- ❑ Reduces the elements of this RDD using the specified **commutative** and **associative binary** operator **f**.


worker  [x334,
x335,
x336,
...,
x665]

f=add



```
val=x1 +  
x2 +  
x3 +  
...  
x1000
```

- ❑ Result return to **driver** program

worker  [x666,
x667,
x668,
...,
x1000]

The Reduce Action

```
def add(x,y):  
    return x+y  
val=myrdd.reduce(add)
```

- ❑ Reduces the elements of this RDD using the specified **commutative** and **associative binary** operator f .

only this is checked!!



- ❑ f is **binary**: takes 2 arguments

- ❑ f is **commutative**:

$$f(x,y) == f(y,x)$$

- ❑ f is **associative**:

$$f(f(x,y),z) == f(x,f(y,z))$$



Examples

```
def add(x,y):                # numbers (addition)
    return x+y               # strings,lists (concatenation, if order
                             #                        does not matter)
                             # arrays, matrices (element-wise addition)
def mult(x,y):               # numbers
    return x*y
                             # numbers, strings
def max(x,y):
    if x>y:
        return x
    else:
        return y

def intersection(x,y):       # collections (lists, sets)
    return Set(x).intersection(y)
```

Other examples: minimum, logical **and**, **or**, **xor**,, bit-wise **and**,**or**,**xor**



Non-Examples

```
def addSquares(x,y):  
    return x**2+y**2
```

commutative but not associative

```
def scaledAdd(x,y):  
    return 2*x + 2*y
```

commutative but not associative

?

associative but not commutative

```
def subtract(x,y):  
    return x - y
```

neither associative nor commutative

```
def ImbalancedAdd(x,y):  
    return 2*x + y
```

neither associative nor commutative



Non-Examples

```
def addSquares(x,y):  
    return x**2+y**2
```

commutative but not associative

```
def scaledAdd(x,y):  
    return 2*x + 2*y
```

commutative but not associative

```
def matrixMult(A,B):  
    return numpy.dot(A,B)
```

associative but not commutative

```
def subtract(x,y):  
    return x - y
```

neither associative nor commutative

```
def ImbalancedAdd(x,y):  
    return 2*x + y
```

neither associative nor commutative



How does reduce work?

```
def add(x,y):  
    return x+y
```

worker



[x1,
x2,
x3,
...,
x333]



partial_sum1=x1+x2+...+x333

```
def add(x,y):  
    return x+y
```

worker



[x334,
x335,
x336,
...,
x665]



partial_sum2=x334+x335+...+x665

```
def add(x,y):  
    return x+y
```

worker




[x666,
x667,
x668,
...,
x1000]



partial_sum3=x666+x335+...+x1000

How does reduce work?

Round 1: Move results to $n/2$ processors



How does reduce work?

Round 1: Combine results

par_sum1+5 par_sum2+6 par_sum3+7 par_sum4+8



worker



worker



worker



worker



worker



worker



worker



worker

How does reduce work?

Round 2: Repeat

par_sum1+5 par_sum2+6 par_sum3+7 par_sum4+8



worker



worker



worker



worker



worker



worker




worker



worker

How does reduce work?

Round 3: Repeat



Serial Reduce Execution Time

Suppose we want to apply binary op (e.g. +) on n elements in `li`

`serial_reduce(li, op)`

Cost of one op (in sec): c

Number of op executions: $n - 1$

Serial execution time: $T_S \approx (n - 1) \cdot c = \Theta(n)$


Parallel Reduce Execution Time

`rdd.reduce(op)`

n elements in rdd

p processors

Stage 1:



Assume rdd is equally partitioned

Each processor has n/p elements

Cost of one op (in sec):

c

Number of op executions:

$$\frac{n}{p} - 1$$

Parallel execution time:

$$T_P^1 \approx \left(\frac{n}{p} - 1\right) \cdot c \approx \frac{nc}{p} = \Theta\left(\frac{n}{p}\right)$$

**Decreases as we
throw more machines
to the problem!**


Parallel Reduce Execution Time

`rdd.reduce(op)`

Stage 2:

n elements in rdd

p processors



Total #of rounds for $p = 2^k$ processors: $k = \log_2 p$

Cost of one op (in sec):

c

Cost to transfer/receive 1 element (in sec):

c'

Each round terminates in:

$c + c'$

**Increases as we
throw more machines
to the problem!**

Parallel execution time:

$$T_P^2 \approx \log_2 p \cdot (c + c') = \Theta(\log_2(p))$$

Speedup

Serial execution time: $T_S \approx (n - 1) \cdot c$

Parallel execution time: $T_P = T_P^1 + T_P^2$
 $\approx \left(\frac{n}{p} - 1\right) \cdot c + \log_2 p \cdot (c + c')$

Speedup:

$$\frac{\text{Best } T_S}{\text{Best } T_P} \approx \frac{(n - 1) \cdot c}{\left(\frac{n}{p} - 1\right)c + \log_2 p(c + c')} = \Theta \left(p \cdot \frac{1}{1 + \frac{p \log_2 p}{n}} \right)$$



Communication Costs

Total #of rounds for $p = 2^k$ processors: $k = \log_2 p$

Total #of messages:

Round 1: $p/2$

Round 2: $p/4$

...

Round $\log p$: 1

$$\begin{aligned} \frac{p}{2} + \frac{p}{4} + \dots + 2 + 1 &= p \sum_{i=1}^{\log_2 p} \frac{1}{2^i} = \\ &= p \cdot \frac{1}{2} \cdot \frac{1 - \frac{1}{2^{\log_2 p}}}{1 - \frac{1}{2}} = p \cdot \left(1 - \frac{1}{p}\right) = p - 1 \end{aligned}$$

total_sum



worker



worker



worker



worker



worker



worker



worker



worker

Communication Costs

- ❑ Reduce **ships aggregated data around**
- ❑ **$\Theta(n)$ messages** for n processors
- ❑ Message size **may increase** with each aggregation!!!
 - ❑ Compare, e.g.,

```
def add(x,y):  
    return x+y
```

applied to **numbers** vs.
applied to **strings** or **lists**



A Map-Reduce Example: Computing the Average

```
rdd=sc.parallelize(range(1000))

total,count =rdd.map(lambda x: (x,1))\
                 .reduce(lambda x,y:
                        (x[0]+y[0],x[1]+y[1]) )

print 1.*total/count
```



Other Useful Actions

```
nums = sc.parallelize([1, 2, 3])
```

```
# Retrieve RDD contents as a local collection  
nums.collect() # => [1, 2, 3]
```

```
# Return K elements  
nums.take(2)    # => [1, 2]
```

```
# Count number of elements  
nums.count()    # => 3
```

**Exercise: how would I
implement these with
a reduce?**



Other Useful Actions

```
nums = sc.parallelize([1, 5, 6, 3, 1, 2 ])
```

```
# Get the N elements from a RDD ordered in ascending order or as  
# specified by the optional key function.
```

```
nums.takeOrdered(3)    # => [1, 1, 2]
```

```
nums.takeOrdered(3, key=lambda(x):-x)    # => [6,5,3]
```

```
# Write elements to a text file
```

```
nums.saveAsTextFile("file.txt")
```

reduce vs. aggregate

- ❑ reduce: **input** and **output** of aggregation are of the same type.

- ❑ aggregate allows them to be different

`aggregate(zeroValue, seqOp, combOp)`

reduce vs. aggregate

❑ reduce: **input** and **output** of aggregation are of the same type.

❑ aggregate allows them to be different

`aggregate(zeroValue, seqOp, combOp)`

Combine pair (OT,IT)

IT: Input Type
OT: Output Type



reduce vs. aggregate

❑ reduce: **input** and **output** of aggregation are of the same type.

❑ aggregate allows them to be different

`aggregate(zeroValue, seqOp, combOp)`

Combine pair (OT,OT)



IT: Input Type
OT: Output Type



reduce vs. aggregate

- ❑ reduce: **input** and **output** of aggregation are of the same type.
- ❑ aggregate allows them to be different

`aggregate(zeroValue, seqOp, combOp)`

Starting OT value
("zero" element)

IT: Input Type
OT: Output Type



Computing an Average through aggregate

```
rdd=sc.parallelize(range(1000))

total,count =rdd.aggregate(
    (0,0),                                     #zero element
    lambda pair,data:                          #seqOp
        (pair[0]+data,pair[1]+1),
    lambda pair1,pair2:                       #combOp
        (pair1[0]+pair2[0],pair1[1]+pair2[1])
)

print 1.*total/count
```



Actions Available for Numerical RDDs only

<code>mean()</code>	Average of the elements
<code>sum()</code>	Total
<code>max()</code>	Maximum value
<code>min()</code>	Minimum value
<code>variance()</code>	Variance of the elements
<code>sampleVariance()</code>	Variance of the elements, computed for a sample
<code>stdev()</code>	Standard deviation
<code>sampleStdev()</code>	Sample standard deviation

Next Lectures

- ❑ Working with key,value pairs
- ❑ Controlling Parallelism
- ❑ Lazy Evaluation & Persistence

