



Northeastern

EECE5698

Parallel Processing for Data Analytics

Lecture 4: Key-Value Pairs & Partitioning

Outline

- ❑ Key-Value Pairs
- ❑ Joins
- ❑ Parallelism & Partitioners



Outline

- Key-Value Pairs

- Joins

- Parallelism & Partitioners



Working with Key-Value Pairs

- ❑ RDDs of **key-value pairs** play a special role in Spark
- ❑ Python: a pair is a **tuple** of **two elements**:

```
pair = (a, b)
pair[0] # => a
pair[1] # => b
```



reduceByKey

```
pets = sc.parallelize([('cat', 1), ('dog', 1), ('cat', 2)])
```


```
pets.reduceByKey(lambda x, y: x + y)  
# => {(cat, 3), (dog, 1)}
```

- ❑ It's a **transform**, not an **action** (produces a new rdd)
- ❑ Presumes that data is in (key,value) pair form
 - ❑ Error will be generated if they are not.
 - ❑ True for all ...ByKey() operations



Example: Word Count

```
lines = sc.textFile("hamlet.txt")
counts = lines.flatMap(lambda line: line.split()) \
               .map(lambda word: (word, 1)) \
               .reduceByKey(lambda x, y: x + y)
```



How Does reduceByKey Work? (simplified)

Each key is mapped to a machine **(actually, partition)**

`target_machine = hash(Key) % num_machines`

using python's builtin `hash()` function

(to,1)
(be,1)

(or,1)
(not,1)

(to,1)
(be,1)

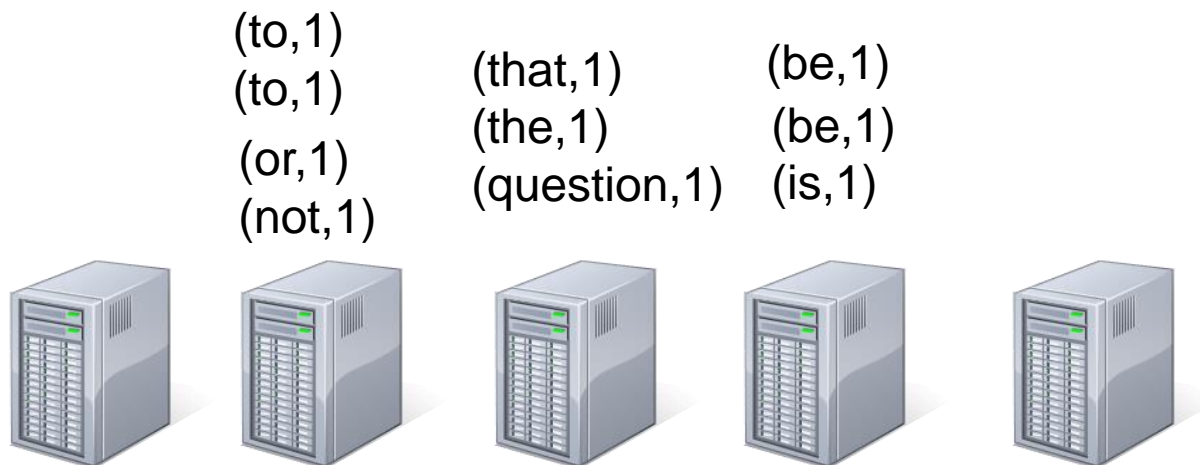
(that,1)
(is,1)

(the,1)
(question,1)




How Does reduceByKey Work? (simplified)

A **shuffle** takes place: key-value pairs are moved appropriate machines, collocating pairs with identical keys



How Does reduceByKey Work? (simplified)

Reduce then applied locally at each machine



Optimizations

- ❑ Values are first combined locally **before the shuffle** takes place
- ❑ **Shuffles avoided** when not needed (partition-awareness, described soon)
- ❑ Big improvement in later versions of Spark: **sort-based shuffle** (we will not cover this)



Additional ByKey Transforms

```
pets = sc.parallelize([('cat', 1), ('dog', 1), ('cat', 2)])
```

```
pets.reduceByKey(lambda x, y: x + y)  
# => {(cat, 3), (dog, 1)}
```

```
pets.groupByKey()  
# => {(cat, Seq(1, 2)), (dog, Seq(1))}
```

```
pets.sortByKey()  
# => {(cat, 1), (cat, 2), (dog, 1)}
```



Transforms Applied on Values only

```
pets = sc.parallelize([('cat', 1), ('dog', 1), ('cat', 2)])
```

```
pets.mapValues(lambda x: x + 1 )  
# => [('cat', 2), ('dog', 2), ('cat', 3)]
```

```
pets.flatMapValues(lambda x: range(x+1))  
# => [('cat', 0), ('cat', 1), ('dog', 0), ('dog', 1), ('cat',  
0), ('cat', 1), ('cat', 2)]
```



Useful PairRDD Transforms/Actions/IO

```
pets = sc.parallelize([('cat', 1), ('dog', 1), ('parrot', 2)])
```

```
pets.values()                # rdd containing values only  
# => [1, 1, 2]
```

```
pets.keys()                  # rdd containing keys only  
# => ['cat', 'dog', 'parrot']
```

```
pets.collectAsMap()          # returns dictionary to driver  
# => {'cat':1, 'dog':1, 'parrot':2}
```

```
allFiles = sc.wholeTextFiles('dir') # loads all files in  
directory 'dir' in a PairRDD, with file names as keys and file  
contents as values
```



combineByKey() -- similar to aggregate()

`combineByKey(createCombiner, mergeValue, mergeCombiners)`

turns a V
into a C

merges a V
into a C

combines two C's
into a single one

Turns an `RDD[(K, V)]` into a result of type
`RDD[(K, C)]`, for a “combined type” C

combineByKey() -- similar to aggregate()

`combineByKey(createCombiner, mergeValue, mergeCombiners)`

turns a V
into a C

merges a V
into a C

combines two C's
into a single one

Turns an `RDD[(K, V)]` into a result of type
`RDD[(K, C)]`, for a “combined type” C

Example: Computing a Per Key Average:

```
keyAvg = words.combineByKey(  
    (lambda x: (x,1)),           #converts val to (val,1)  
    (lambda x, y: (x[0] + y, x[1] + 1)), #adds val to running tally  
    (lambda x, y: (x[0] + y[0], x[1] + y[1])) #merges two tallies  
) .mapValues(lambda (val,count):1.*val/count)
```

Outline

□ Key-Value Pairs

□ **Joins**

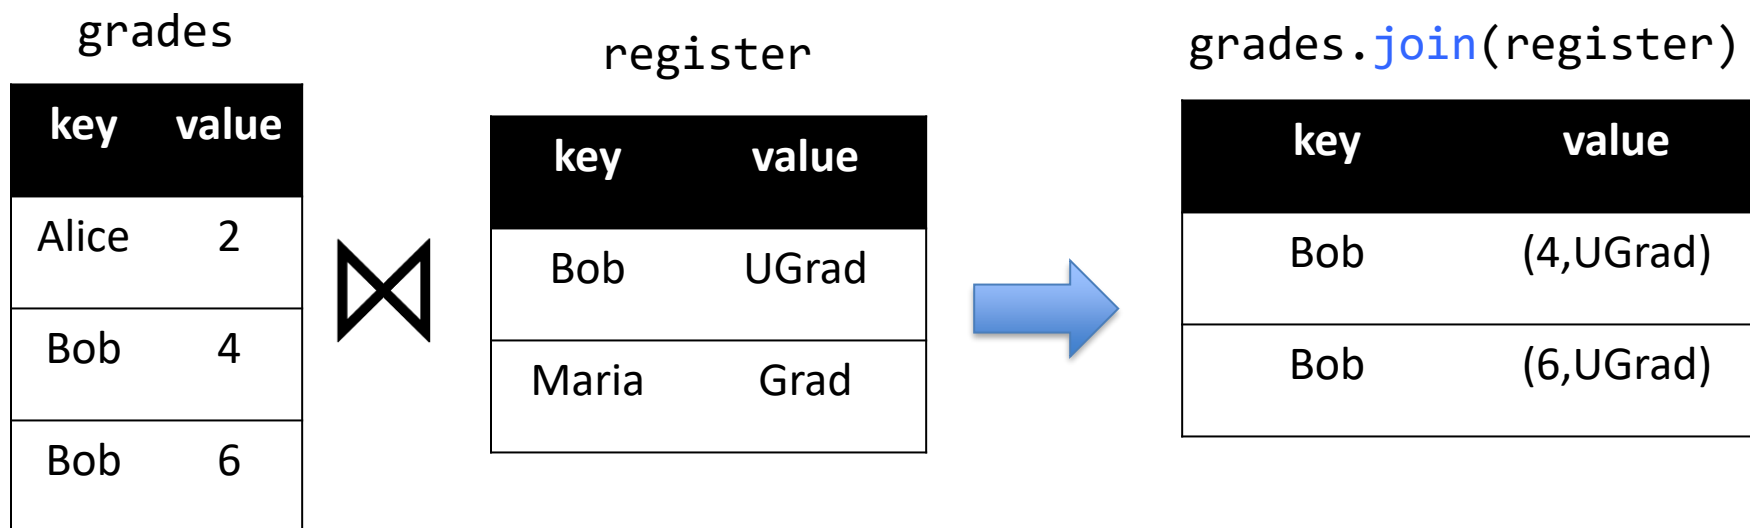
□ Parallelism & Partitioners



Joins

```
grades = sc.parallelize([('Alice', 2), ('Bob', 4), ('Bob', 6)])  
register = sc.parallelize([('Bob', 'UGrad'), ('Maria', 'Grad')])
```

```
grades.join(register)          # Perform an inner join between two RDDs  
# => [('Bob', (4, 'UGrad')), (Bob, (6, 'UGrad'))]
```



If (key, val1) in rdd and (key, val2) in other, join contains (key, (val1, val2))

How Does join Work? (simplified)

Joins require **shuffling**

rdd:grades

rdd:register

(Maria, Grad) (Bob,4)
(Alice,2) (Bob,6) (Bob,Ugrad)



How Does join Work? (simplified)

Joins require **shuffling**

rdd:grades

rdd:register

(Maria, Grad) (Bob, Ugrad)
(Alice, 2) (Bob, 4)
(Bob, 6)



How Does join Work? (simplified)

Joins require **shuffling**

rdd:grades
rdd:register

Shuffles avoided when not necessary
through **partition-awareness**

(Bob,(4,Ugrad))
(Bob,(6,Ugrad))



Other Transformations On Pairs of RDDs

```
rdd = sc.parallelize([(1, 2), (3, 4), (3, 6)])  
other = sc.parallelize([(3,9),(5,8)])
```

```
rdd.join(other) # Perform an inner join  
# => [(3, (4, 9)), (3, (6, 9))]
```

```
rdd.leftOuterJoin(other) # Perform a left outer join  
# => [(1, (2, None)), (3, (4, 9)), (3, (6, 9))]
```

```
rdd.rightOuterJoin(other) # Perform a right outer join  
# => [(3, (4, 9)), (3, (6, 9)), (5, (None, 8))]
```

```
rdd.subtractByKey(other) # Remove elements with key in other  
# => [(1,2)]
```

```
rdd.cogroup(other) # Group data sharing the same key together  
# => [(1, ([2], [])), (3, ([4, 6], [9])), (5, ([], [8]))]
```



Outline

□ Key-Value Pairs

□ Joins

□ **Parallelism & Partitioners**



Controlling the Level of Parallelism

All the pair RDD operations, and some of the non-pair operations, take an optional second parameter for **number of partitions**

```
words.reduceByKey(lambda x, y: x + y, 5)
```



```
words.groupByKey(5)
```



```
visits.join(pageViews, 5)
```






This can be used to control the **level of parallelism**

What Are Partitions?

□ RDDs are internally split into **partitions**

```
myrdd=sc.textFile("WarAndPeace.txt",3)
```



What Are Partitions?

partitions < #machines

part0000



part0001



part0002



What Are Partitions?

partitions > #machines

part0005
part0000



part0004
part0009



part0006
part0001



part0003
part0002




part0007
part0008



Number of Partitions Controls Parallelism

```
rdd.map(lambda x: x+1)
```

- ❑ Map executed **serially** in each partition
- ❑ If machine stores **k** partitions, and has **n>k** processors, partition evaluations **executed in parallel**



How Many Partitions Should One Use?

- ❑ m workers with k processors each

"ideal" #partitions = $m k$

- ❑ **Fewer partitions:** not exploiting full parallelism in this operation
- ❑ **More partitions:** No advantage in speedup; in practice, there may be advantage in memory usage (each cpu dealing with smaller partition, spark less likely to crash/hang)



Mapping Data to Partitions

Each key is mapped to a machine  **(actually, partition)**

`target_machine = hash(Key) % num_partitions`

using python's builtin `hash()` function

(to,1)
(be,1)



(or,1)
(not,1)



(to,1)
(be,1)



(that,1)
(is,1)

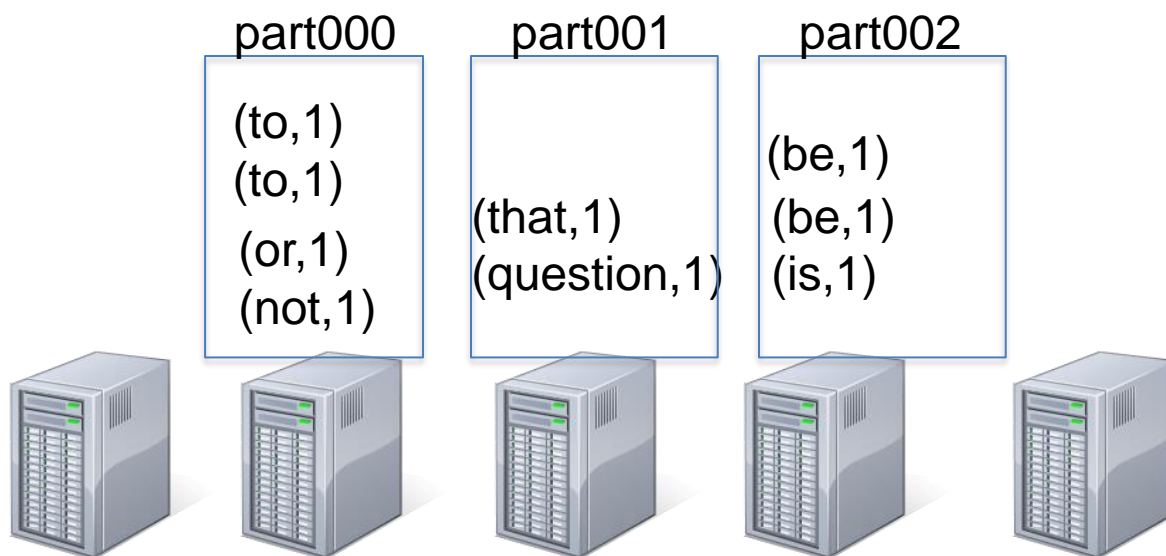


(the,1)
(question,1)



How Does reduceByKey Work? (simplified)

A **shuffle** takes place: key-value pairs are moved appropriate machines, collocating pairs with identical keys



Re-Partitioning a key-value pair RDD

```
words.partitionBy(numPartitions,partitionFunc=hash)
```

```
words.partitionBy(5) # create 5 partitions with default hash
```

```
words.partitionBy(5,partitionFunc=lambda x:hash(x)+10)
```

```
# create 5 partitions with user-defined hash
```

```
somerdd.repartition(5)
```

```
# partition a non key-value pair rdd
```



Partition-Awareness

```
lines = sc.textFile("hamlet.txt")
counts = lines.flatMap(lambda line: line.split()) \
               .map(lambda word: (word, 1)) \
               .reduceByKey(lambda x, y: x + y, 100)
```

rdd **not**
partitioned,
reduceByKey
requires a
shuffle

```
lines = sc.textFile("hamlet.txt")
```

```
wordsPartitioned= lines.flatMap(lambda line: line.split()) \
                     .map(lambda word: (word, 1)) \
                     .partitionBy(100)
```

rdd **is partitioned**,
reduceByKey does
not shuffle!

```
counts = wordsPartitioned.reduceByKey(lambda x, y: x + y)
```


How Does This Work?

- ❑ key-to partition map fully defined by
 - ❑ NumPartitions
 - ❑ PartitionFunc (default:hash)
- ❑ When rdd is shuffled by partitionBy, these are stored in private variables
- ❑ When reduceByKey() is called next, it checks to see if these fields are set; if so, it skips shuffling



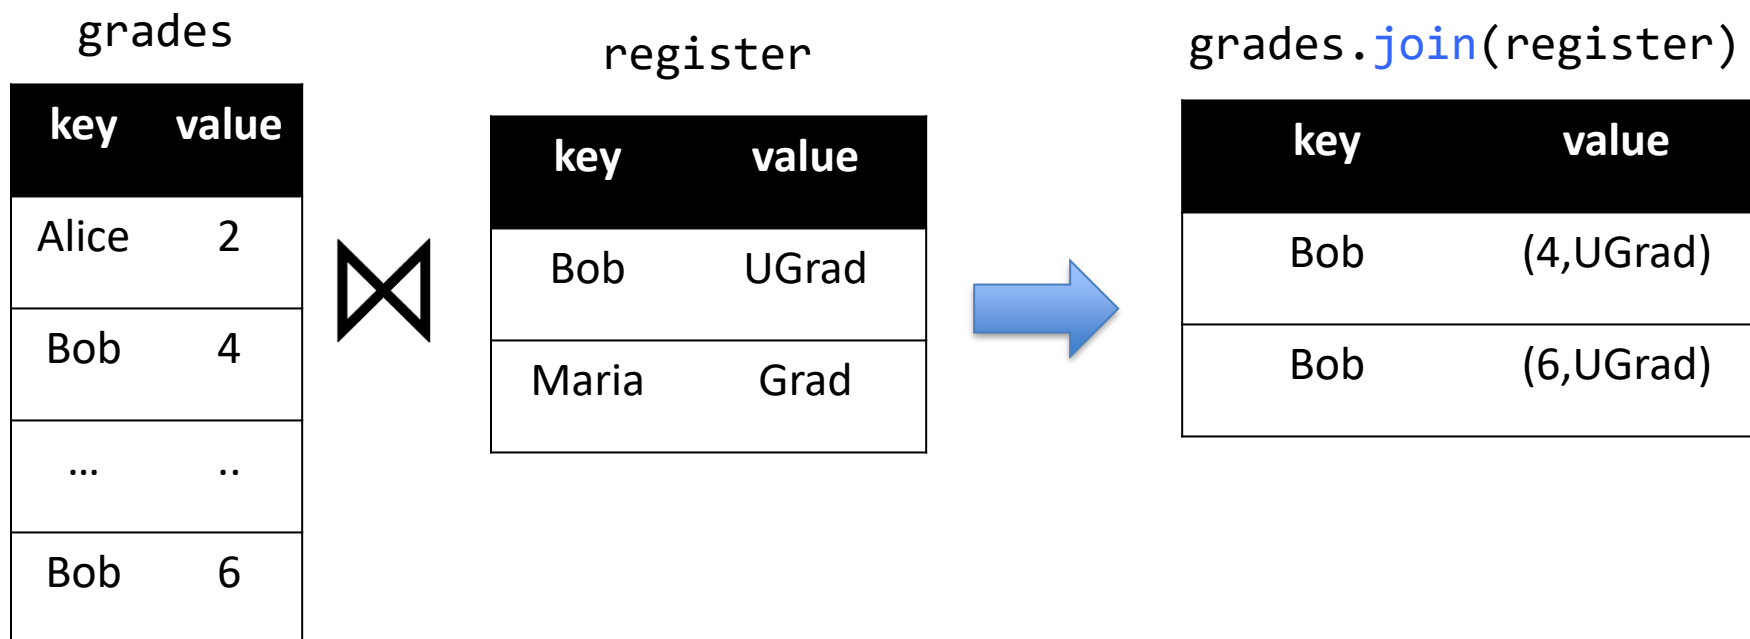
Partition-Awareness & Joins

- ❑ key-to partition map fully defined by
 - ❑ NumPartitions
 - ❑ PartitionFunct (default:hash)
- ❑ Join of rdds with **the same partitioning information** do not require a shuffle!!!!
 - ❑ keys are already on the same machines
- ❑ Useful when doing repeated joins on large rdd



Partition-Awareness & Joins

```
grades = sc.parallelize([('Alice', 2), ('Bob', 4), ..., ('Bob', 6)])\  
    .partitionBy(100).cache() #Perform a shuffle, sets partition info  
...  
register = sc.parallelize([('Bob','UGrad'),('Maria','Grad')])\  
    .partitionBy(100) #Perform a tiny shuffle, sets partition info  
  
grades.join(register)          # NO SHUFFLING REQUIRED
```



Operations That Create/Preserve Partitioner

Create & Preserve

- ❑ `cogroup()`, `groupWith()`
- ❑ `join()`, `leftOuterJoin()`, `rightOuterJoin()`,
- ❑ `groupByKey()`, `reduceByKey()`, `combineByKey()`
- ❑ `partitionBy()`, `sortByKey()`

Preserve (if parent has partitioner)

- ❑ `mapValues()`, `flatMapValues()`
- ❑ `filter()`

Beware of non-key value operations!

- ❑ A map will **remove partitioning info**, even if it does not alter keys.

